


CSE30 Final Review

This review doc summarizes essential concepts covered in lectures. For detailed explanation, please also kindly refer to in class examples. Created by Yilin.

Number representation:

1. Unsigned binary number - decimal $101_2 = 3_{10}$
 - a. Signed
 - i. 1st bit to represent +/-, rest bits represent **magnitude**.
 - ii. $101_2 = -1$;
2. 2's complement
 - a. Inverting all bits and add 1(negative numbers)
 - i. $-2_{10} = 0010$ in magnitude
 - ii. Flip bits => 1101
 - iii. Add 1 => 1110
 - iv. 1001_2 flip bits => 0110₂
 - v. Add 1 => 0111₂
 - vi. $0111_2 = 7_{10}$, signed is negative => -7_{10}
3. Hex binary
4. Hex decimal
5. Float
 - a. How to represent numbers like $\frac{1}{2}$, 0.001, 6.023×10^{23}
 - b. IEEE standard floating point




- c. **Always have a 1 in the beginning of the mantissa(implicit existence)**
- d. **Exponent = exponent bits - 127(bias)**
- e. **zero:** exponent all 0s, mantissa all 0s.

Bitwise operation

1. and/ or/ xor
 - a. And: both inputs are 1, output 1, otherwise output 0
 - b. Or: both inputs are 0, output 0; otherwise output 1
 - c. Xor: exactly one of the input is 1, output 1; otherwise output 0
2. Masking - use add to clear bits

Combinatorial logic

1. Show a circuit, and fill the truth table
 - a. AND operator: returns 1 when both of the inputs are one
 - b. OR operator: returns 1 when either one of the inputs is one
 - c. XOR operator: returns 1 only when one of the inputs is one
2. Full adder:



3. full adder.

ASCII & C strings

1. Null-terminated strings
 - a. "Hello" = 'H' 'e' 'l' 'l' 'o' '\0'
 - b. Length = 5 (strlen()的结果)
 - c. Char str[4] = "four"; no null terminator(garbage data)
 - d. Char str[5] = "four"; null terminator
 - e. Char str[3] = "four"; **warning!**
2. Operations on strings
 - a. Strdup - copy
 - b. Strlen - return the length of a string. # of characters before encounter '\0'.
 - c. Changing characters

Undefined Behavior:

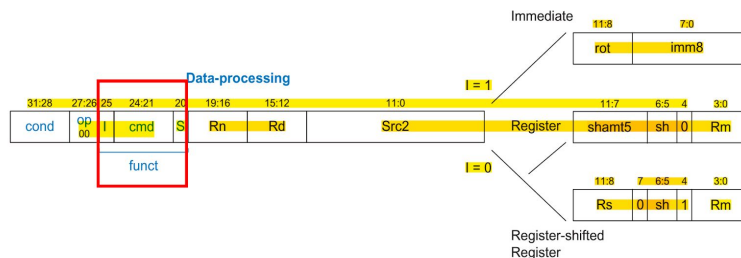
- `char* my_str = "asdf";`

- Do not know where it is stored due to undefined behavior
- To solve this, we use `strdup`, which copies "asdf" and stores it in **heap**
- Examples
 - Memory Leaks
 - Incorrect execution (crashes, incorrect results)
 - Accidentally correct execution
 - I.e. Garbage values just so happen to line up with desired values
 - Reading/writing memory out of bounds
 - I.e. allocating room for 3 words (12 bytes) on the stack, then trying to access the 13th byte
- `arr[n]` means: `*(arr + n)`
- `int size = 10;`
`int * some_num = malloc(sizeof(int) * size);`
`Some_num[12] = 1;`
- **Cause undefined behavior: could have changed the value stored after `some_num`.**

Machine code

1. Encoding of register

cond	Mnemonic	Name	CondEx
0000	EQ	Equal	Z
0001	NE	Not equal	\bar{Z}
0010	CSHS	Carry set / unsigned higher or same	C
0011	CCLO	Carry clear / unsigned lower	\bar{C}
0100	MI	Minus / negative	N
0101	PL	Plus / positive or zero	\bar{N}
0110	VS	Overflow / overflow set	V



Basic Assembly functions

1. Mov
2. Add
3. Sub
4. Ldr
5. Lsl - logic shift left:
 - a. `0x80000000 -> 0x0000 0000`
 - b. `N = 0, Z = 1, C = 1, V = 0`
6. ...

CPSR

1. `0xFFFFFFFF (-1) + 0x00000001 (1) = 0x00000000(-0)`
 - a. Does not set N bit because result is not negative
 - b. Does set Z bit because result is **zero**
 - c. Does set C bit because there is a carry
 - d. Does not set V bit because it's only set when two + -> -
 - i. eg: `0100 0000 + 0100 0000 -> 1000 0000`

2. $0xFFFFFFFF (-1) + 0x00000002 (2) = 0x00000001 (1)$
 - a. Carry bit is set
3. Running commands based on CPSR

cond	Mnemonic	Name	CondEx
0000	EQ	Equal	Z
0001	NE	Not equal	\bar{Z}
0010	CS/HS	Carry set / unsigned higher or same	C
0011	CC/LO	Carry clear / unsigned lower	\bar{C}
0100	MI	Minus / negative	N
0101	PL	Plus / positive or zero	\bar{N}
0110	VS	Overflow / overflow set	V

Labels and Loops

1. Program counter
 - a. Where we are in a program (which **instruction**)
 - b. Increase by 4 because every instruction is 4 bytes
 - c. If ask for an instruction, pc is always 8 bytes ahead
2. B label
 - a. Branch
 - b.
3. Loops
4. If statement

Memory instructions

1. Loading and storing bytes
2. Loading and storing words
3. Pushing and popping

Calling

1. bl command

- a. r0 - r3 are often overwritten in the process of calling a function. Hence, the caller must save those registers if it depends on any of its own arguments (original values of r0 - r3) after a function returns.

Preserved	Nonpreserved
Saved registers: R4–R11	Temporary register: R12
Stack pointer: SP (R13)	Argument registers: R0–R3
Return address: LR (R14)	Current Program Status Register
Stack above the stack pointer	Stack below the stack pointer

2. Stack pointer's role

- a. **Close to where the command-line arguments are stored when the program starts**

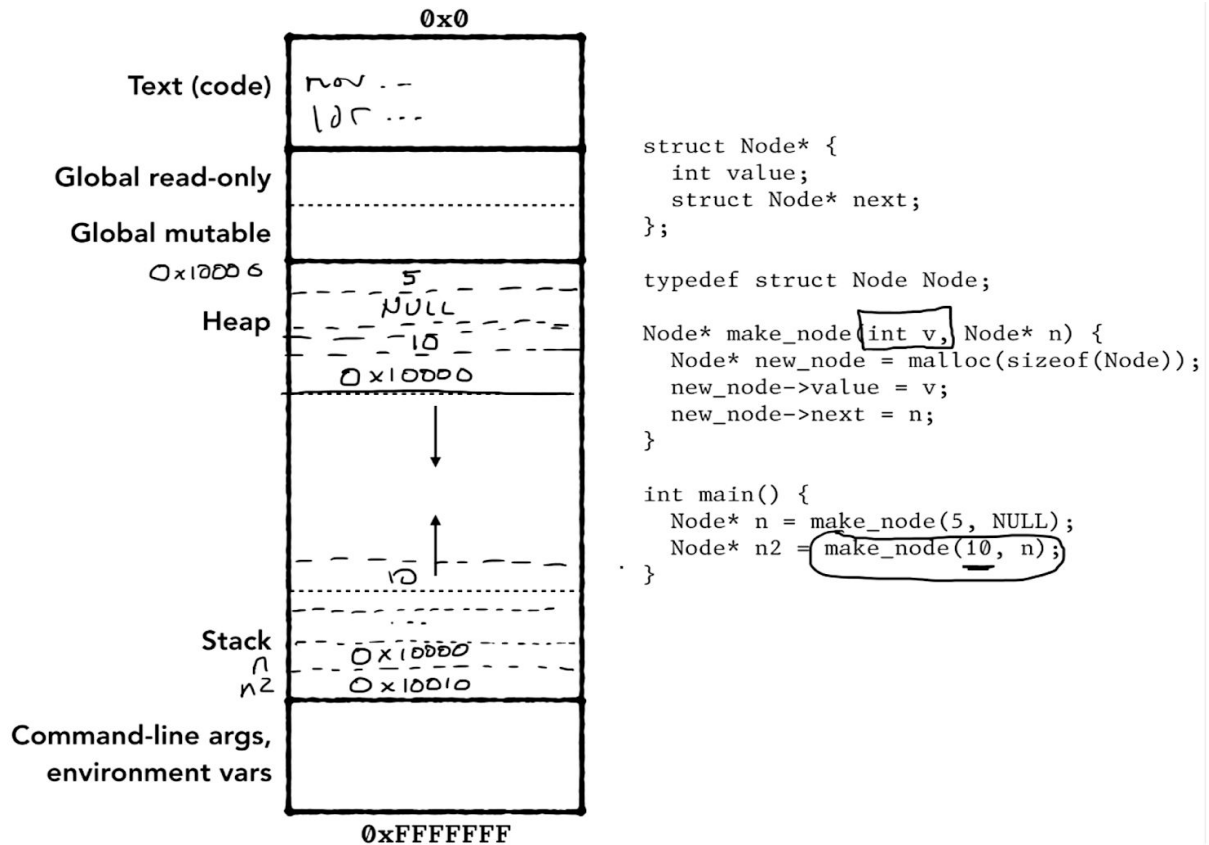
3. Callee-save and caller-save

- a. Caller must preserve these registers: r0, r1, r2, r3 and r12 (caller save rule)
 - i. Stack below the sp
- b. Callee must preserve these registers: r4 - r11, sp, lr (callee save rule)
 - i. Stack above the sp

4. Passing in arguments in registers

5. Restoring before returning

- a. Return from a function: `MOV pc, lr`
- b. `PUSH {lr}` at the start of a non-leaf function + `POP {pc}` at the end
 - This is done to maintain the stack frames of functions
 - Link register is stored on the stack to return to the caller after the callee terminates
- Recursive function calls



Heap memory

- Using malloc for heap-allocated arrays (has todo with sp)
 - Argument: number of **bytes**
 - Returns an **address**
- Using malloc for struct data
- Appropriate use of free()
- global variable is stored in global mutable block**

Stack Memory

- Where stack memory should go for struct and array
- Correct copying and referencing
 - Stack copying: when the pass in arguments are not pointers, copy to stack
 - `f(Pointer p) {p.x = 22;}` - calling f **does not** change the value of p in main
 - `f(Pointer* p) {p.x = 22;}` - calling f **does** change the value of p in main
 - Reminder: pass in an array is same as pass in pointer**
 - In java, it's more like the code with pointer because java heap allocates all its objects
- How & gives the address


The type of x is...	The compiler generates code to...	Example
primitive (int, char)	Pass (copy) directly to callee	<code>int x = 10; f(x); // 10 passed in r0</code>
pointer	Pass (copy) directly to callee; copies an address	<code>int* x = malloc(sizeof(int)); f(x); // address returned from // malloc passed in r0</code>
array	Pass address of array directly to callee	<code>char cs[] = "abcd"; f(cs); // address for start // of cs passed in r0</code>
struct	Copy struct contents to callee	<code>struct Point p = {1, 4}; f(p); // 1, 4 copied to stack // frame for f to use</code>

Stack Layout

- Grows in decreasing address order

- `char s1[] = "Happy";`
`char s2[] = "Tuesday";`
"Happy" is stored at higher address and "Tuesday" is stored at lower address
- `char s1[] = "Happy";`
`int fs = 0xffff;`
`char s2[] = "Tuesday";`
`int ds = 0xdddd;`

Because gcc puts local first, and arrays afterwards



gcc结果

Malloc and free implementation

- `char s[4] = "hello!";` `sizeof(s)` -> 4 instead of 7
- `char* s = strdup("hello!");` `sizeof(s)` -> 4 instead of 7 because only one word

3. `char* s = strdup("hello!"); sizeof(*s) -> 1` because `*s` returns a `char` (type)
4. `char s[10] = "hello!"; sizeof(s) -> 10`
5. `Struct A{ char s[10]; char* name;}; sizeof(A) -> 16` because multiple of 4 (padding) (largest field)
6. `Struct A{ char s[10]; }; sizeof(A) -> 10` because treated as char array
7. `Struct A{ char s1[5];char s2[9]; }; sizeof(A) -> 14` because rounded to the biggest size
8. `Struct A{ short s1[5];char s2[9]; }; sizeof(A) -> 20` because rounded to the biggest size
9. `Struct A{ char s1; char s2; int i; }; sizeof(A) -> 8` because char can be fitted in
10. `Struct A{ char s1; int i; char s2; }; sizeof(A) -> 12` because char is separated
11. If it comes to the case where one struct has a struct field, only consider the largest field in those struct.
 - a. `Struct A {char c1; int i;}; struct B {char c2; struct A;}; sizeof(B) -> 12`; c2 takes 4 bytes of space because the largest field is int in these two structs.
 - b. `Struct A {char c1; int i;}; struct B {struct A; char c2;}; sizeof(B) -> 12`, same thing
12. Signature:
 - a. `void* malloc(size_t size);`
`void* free(void* ptr);`
 - b. **PA6**

Caching

1. The index in a direct-mapped cache, how to evict/ replace data
2. CPU:
 - a. Registers
3. Random-access memory
 - a. Connected with CPU in some physical distance
 - b. Accessing the data in RAM is 100 times longer than accessing that in registers
 - i. Instructions such as: LDR, STR
4. **Cache:** for recently used memory
 - a. Not very large (in class 8 entries)
 - b. Fetch address:
 - i. 0x0001 0004
Last 2 bytes: 0 4
0000 0100
Last 5 bits - 1st 3: the index in cache; last 2: byte offset
 - ii. **Byte offset:**
Which part of the word u want
 - c. **Representation:**
 - i. 1st bit to indicate use
 - ii. 27 bits for tag
 - iii. 32 bits for data
 - d. How to get data out?
 - i. Go to index first
 - ii. Compare two tags, match? Return data : update the cache

e. Str instruction:

- i. **Write-through:** Write to cache, and to memory(update cache & memory)
- ii. **Write-back: Dirty bit: (set dirty bit 1st)**
 - 1. **When in the future, encountering an address with \neq tag, write into memory. (waits to save to memory)**

Virtual Memory

- 1. Sharing the same physical memory
 - a. In process, same size of memory used in physical memory, but not necessary exactly the same address.
 - b. 1st 5 bytes: **page (different size) PAGE TABLE**
 - i. Page could be different in physical memory
 - ii. Page table \approx 4MB
 - iii. Each entry gets an index
 - iv. 1 bit indicates use or not
 - 1. If there is mapping, set to 1
 - v. If page table is full, **data overlapping**
 - 1. Solution: **mmap (asking OS for more memory)**
 - vi. **Access address that is not mapped in the page table, segmentation fault. (OS has not blessed it....)**
 - vii. **If we use more and more memory, OS automatically allocate pages until segment fault.**
 - viii. What if the physical memory is **full**
 - 1. MMU(memory management unit in CPU)
Only knows one page table at one time
Switch page tables so that convert to the right page table to run the process.
 - 2. OS -
 - a. telling MMU which process is running so which page table to use
 - b. Track used space and swap if needed
 - ix. **SWAP???(virtual memory \rightarrow hard disk, if physical space not enough)**
 - c. 3 left bytes: **offset**